Step-by-step guide for deploying a Streamlit application on an AWS EC2 instance

Aiko Hassett · Follow
5 min read · Jul 13, 2024

Launch a new EC2 instance

From your AWS CLI, go into your **EC2 Dashboard** and select **Launch instances.**

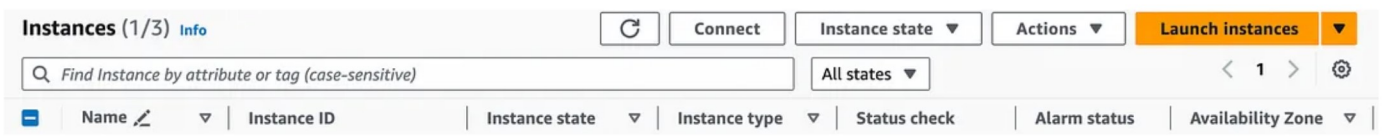

Snippet of the EC2 dashboard

. . .

Configure your EC2 instance

When configuring your **instance type**, pay attention to how big your application is (including all the dependencies that come with running it). For example, a *t2.micro* might not be the best choice if you're looking to run compute-intensive workloads directly from your EC2 instance. For me, my web application ran on a handful of dependencies and API calls, so a *t2.micro* seemed like an adequate option.

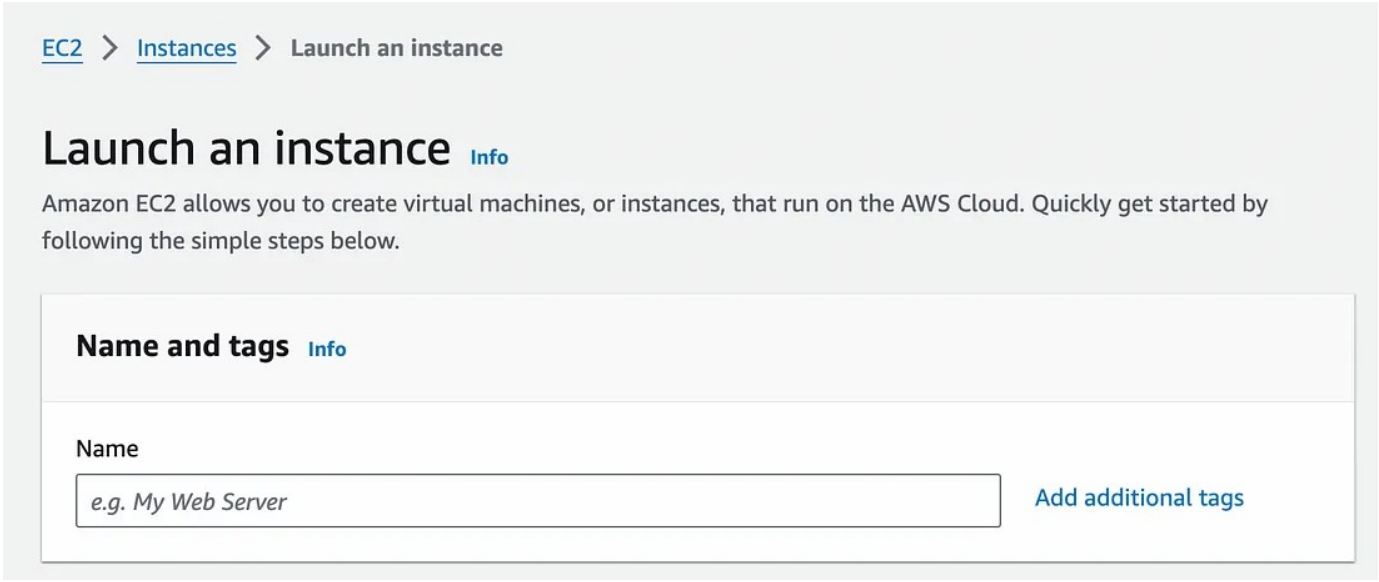

Snippet A — Configure your EC2 instance

Note: At the time of writing this article, t2.micro supports 1GB vCPU and 1GB of memory. Keep in mind, you won't be able to allocate the full 1GB of memory on the t2.micro instance to your web application because some of that will be reserved for use by the operating system. If you are unsure if a t2.micro fits your use case, there is no harm in trying it out and scaling up the instance if it runs out of compute.

Set up a key pair for your EC2 instance

You can use an existing one, or let AWS create a new one for you. There is also an option to proceed without one, but it is highly recommended you set one up to secure your connection to your instance.

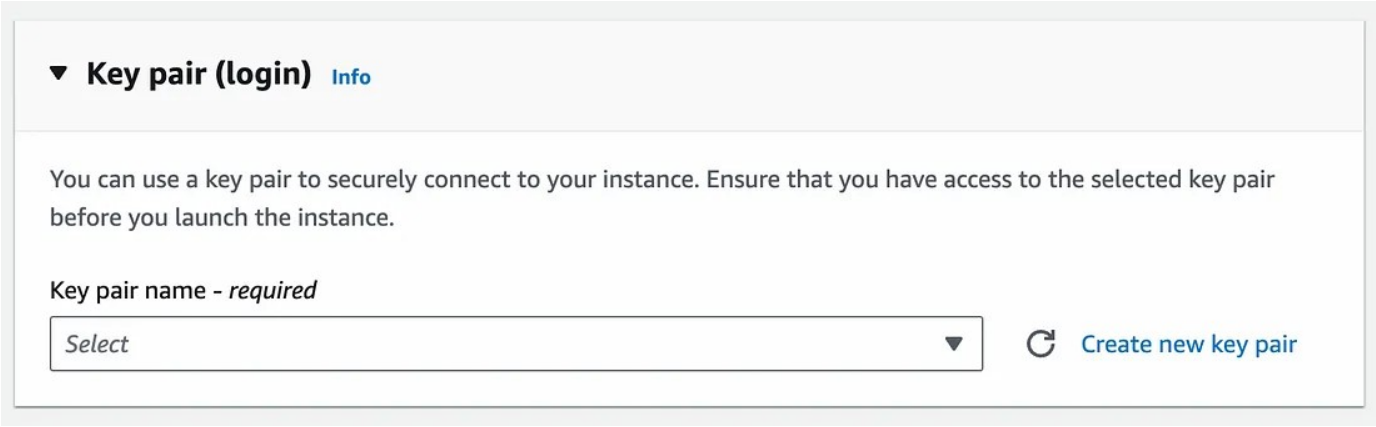

Snippet B— Configure your EC2 instance

Configure your network settings

Make sure you check the appropriate options if you are looking to expose your web application over HTTP (port 80) and/or HTTPS (port 443).

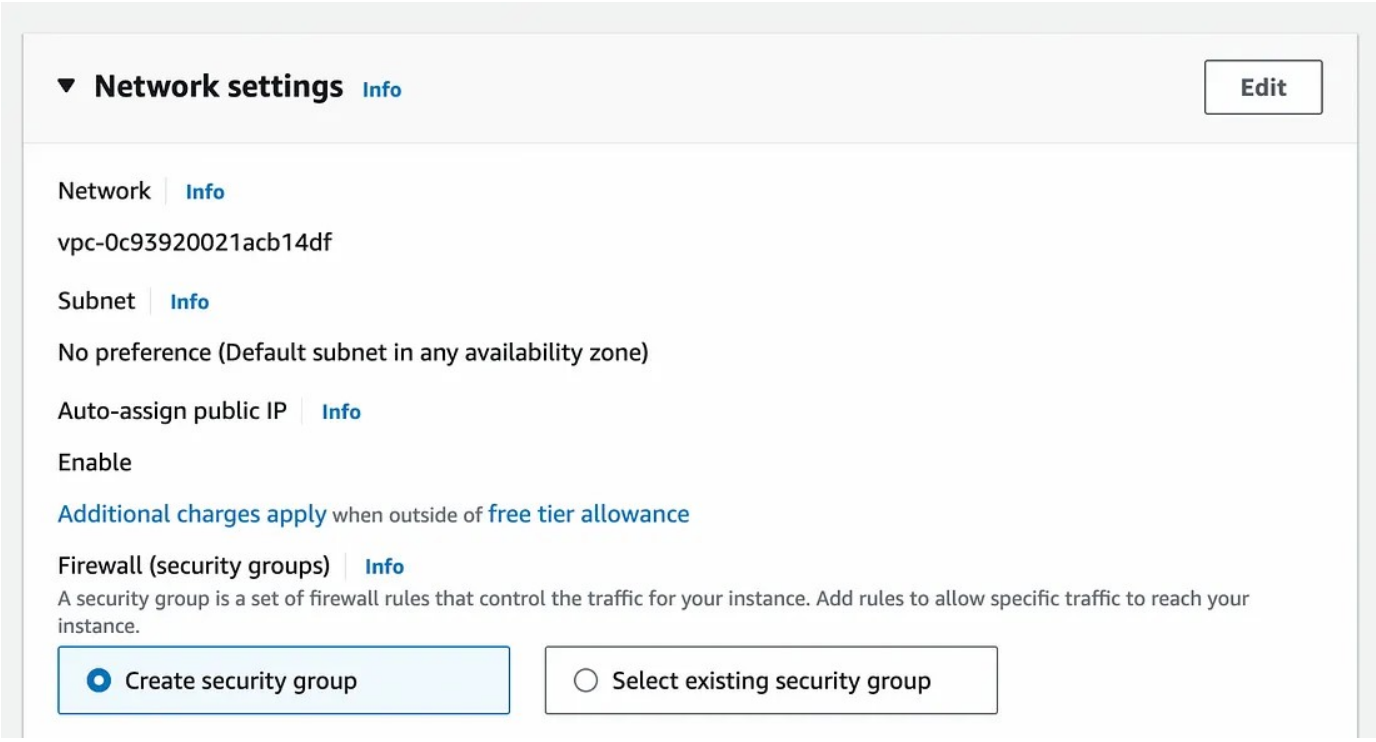

Snippet C— Configure your EC2 instance

. . .

Connect to your EC2 instance

Once your are done configuring your instance details, finalize your changes via **launch instance.** This should be a fairly quick process that can take

anywhere from 30 seconds to a few minutes. You should be redirected to your **Instances** dashboard where you can see the **instance state** update.

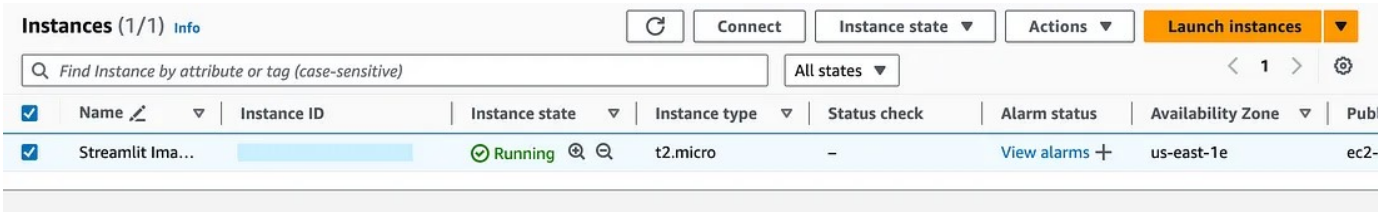

Snapshot D — Connect your EC2 instance

Once the **instance state** shows as **running**, that is your sign to select the checkbox next to your instance and click on **connect.**

From **EC2 Instance Connect**, leave the configuration as is and click **connect.**

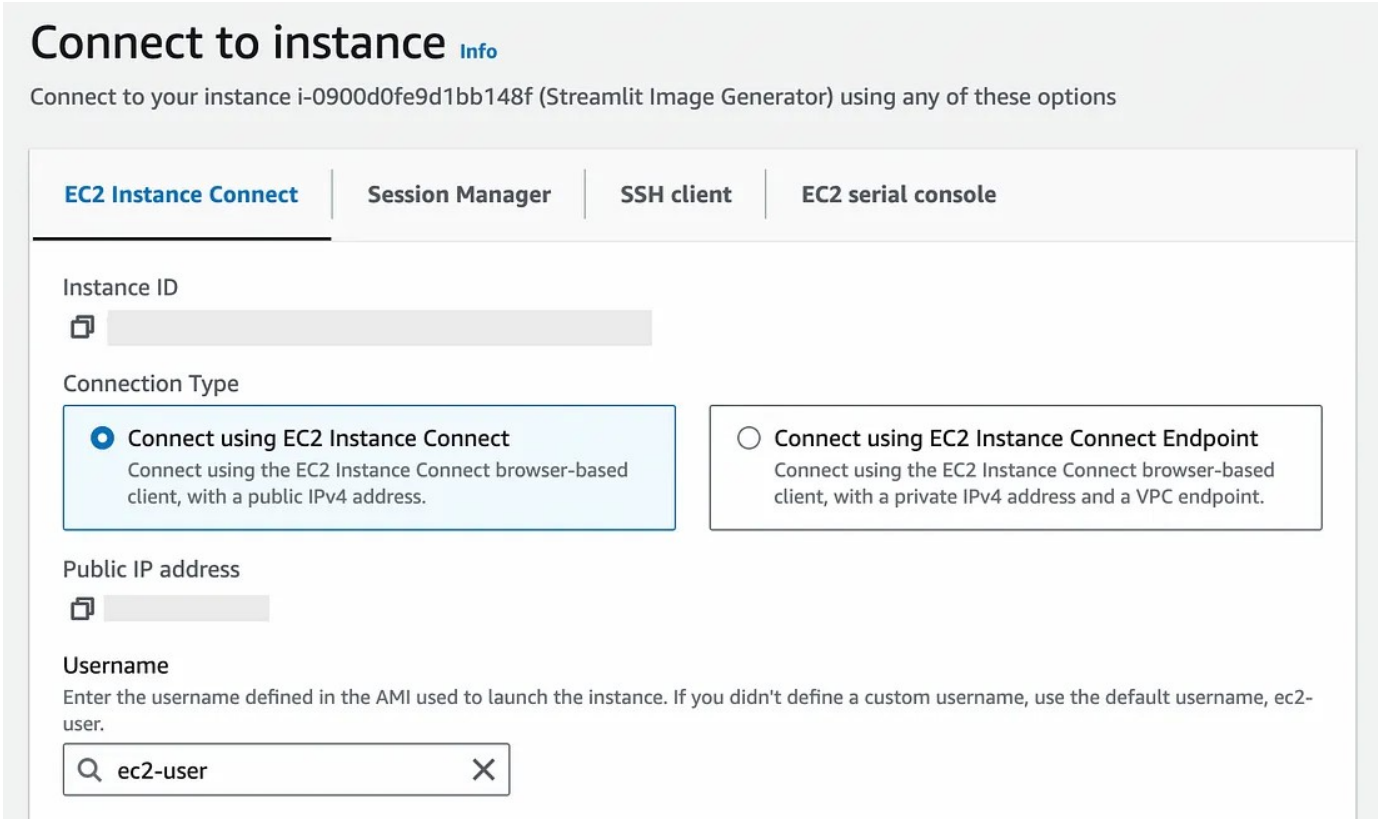

Snapshot E — Connect your EC2 instance

. . .

Setting up your EC2 instance

Once the CLI loads, type in *sudo su* to invoke **root privileges.** This will allow

you to install all your code and it's dependencies under root without any mistakes. This is important because we will need to configure our user data scripts later to run on these files, and user data runs as a root user.

Package installation

Remember, the VM comes as a blank slate, so you will need to reinstall all your dependencies from scratch. Pip and Git are your standard two, you may need to install Docker as well if you are contanerizing your application beforehand.

```
# install pip (package manager)
wget https://bootstrap.pypa.io/get-pip.py
python3 ./get-pip.py

# install git (for those using github/gitlab or similar)
sudo yum install git

# install and build docker (if you are looking to contanerize application)
sudo yum install docker -y
sudo service docker start
sudo docker build -t <my-username/my-image> .
sudo docker login
sudo docker push <my-username/my-image>

# pip install your packages individually
pip install <package-name>

# ...or via requirements.txt
pip install -r requirements.txt
```

Note: EC2 will not allow you to connect to your repository via username and password and will essentially require an SSH connection.

Set up an SSH connection between your EC2 instance and your code repository. The installation process for GitHub can be found here. Clone your repository via SSH and confirm that the operation was successful.

Optional: Create an IAM Role for your EC2 instance

If you are utilizing AWS resources in your code, create an **IAM Role** for your

EC2 instance that gives it access permission to those resources.

Once completed, attach the IAM Role to the instance by going back to the **Instances** dashboard > select your instance and going into **Action > Security > Modify IAM Role** from the dropdown menu on the top right.

Optional: Add your run commands in EC2 User Data

If you are looking to save your instance as an AMI and start up additional instances that automatically run your script at initial launch, this is a section for you.

From your EC2 dashboard, go to **Actions > Instance Settings > Edit User Data.**

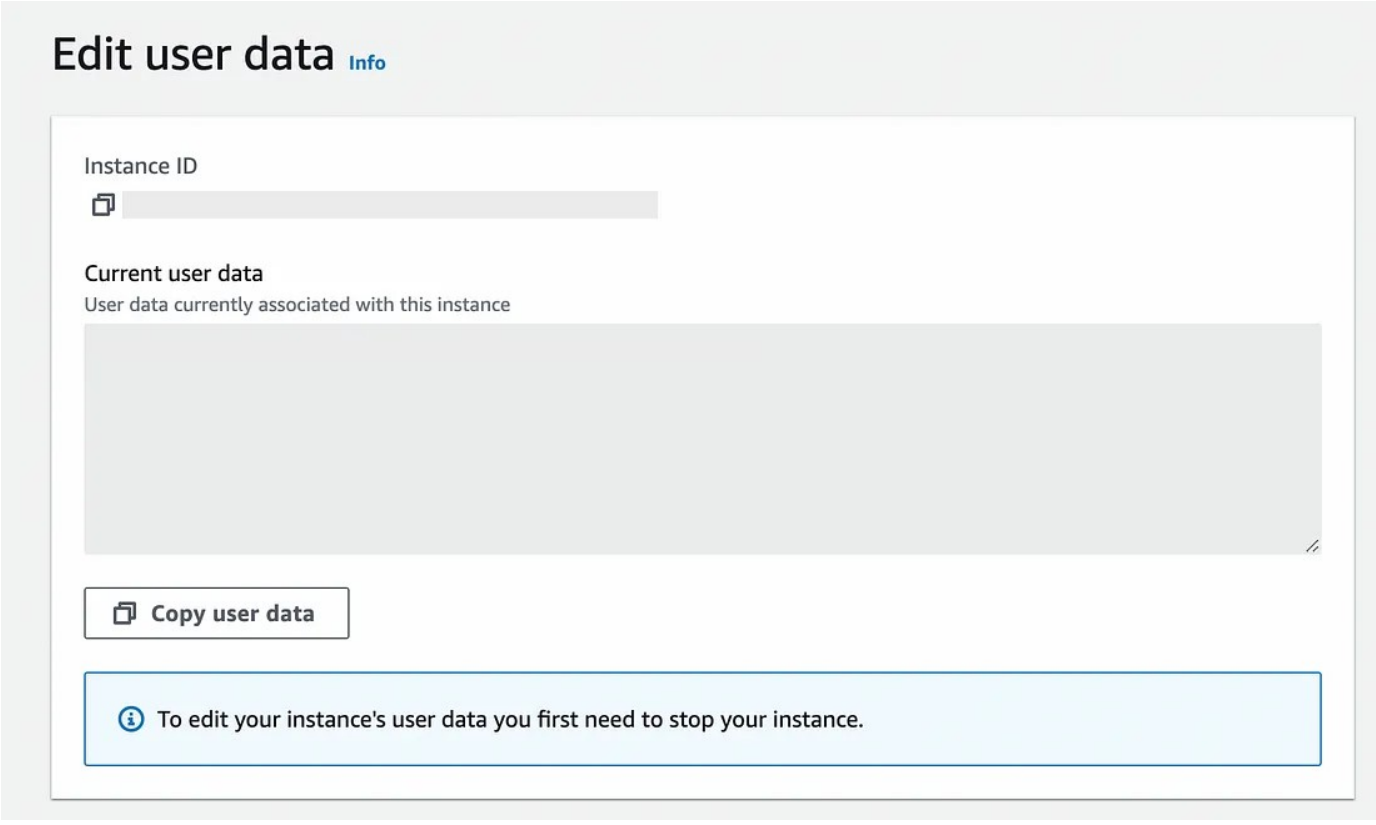

Snapshot F — Add your run commands in EC2 User Data

What you include in here can automate task flows at launch time. As an example, this is what I entered to launch my Streamlit application.

```
streamlit run <path/to/streamlit/><streamlit-app.py>
```

. . .

Confirm everything is working

Go to your EC2 dashboard and copy your **Public IPv4 DNS** address into a new tab. If the URL loads your application, you are good to go. If it stalls, check your network security group configurations, make sure you are mapping to the right port and try again.

Hope this tutorial is helpful to you and good luck on your AWS journey!

Streamlit AWS Ec2 Cloud Computing Walkthrough

Written by Aiko Hassett

Follow

0 Followers

Hello. I'm passionate about the cloud and love sharing my knowledge on Medium.

Recommended from Medium

Manyi

Hosting Streamlit App on AWS EC2

If you have created a Streamlit app and like to deploy it on AWS, for your own use or for...

Aug 3 6

Krishan Walia

Master Machine Learning Pipeline Development | MlOps Project-1 |...

Series of articles with complete explanation, that can help you master machine learning...

Sep 15 2 1

Lists

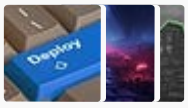

Predictive Modeling w/ Python

20 stories · 1622 saves

Natural Language Processing

1780 stories · 1386 saves

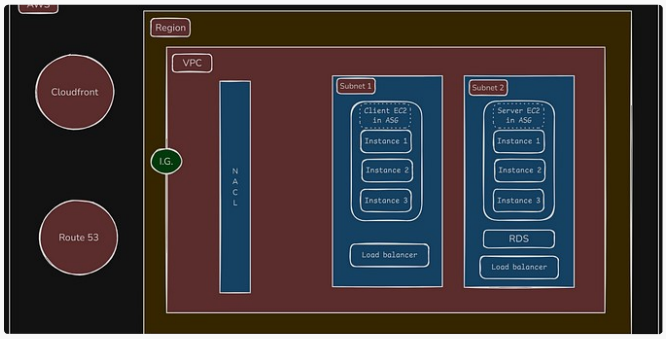

pritesh

Learning Streamlit part—4

33. Custom Theming in Streamlit

5d ago

Samriddha Chattopadhyay in AWS Tip

AWS networking

Networking in AWS is a really important concept to understand as a beginner in...

Oct 10 35 3

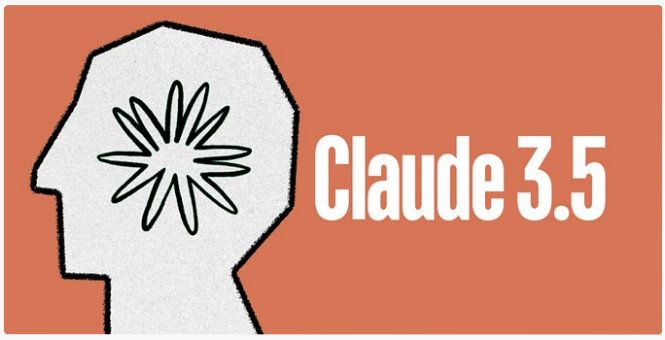

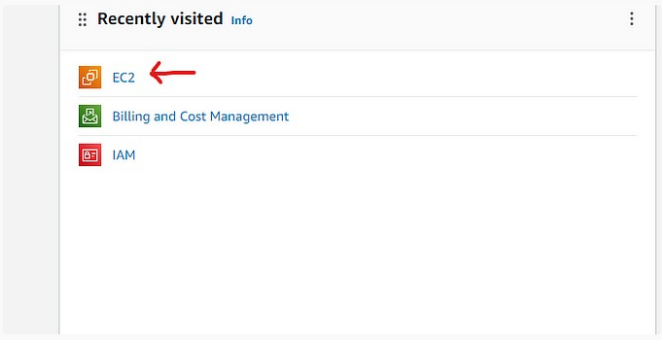

Manpreet Singh in AI Advances

AI are using computers like humans: Claude 3.5 new update

Hey everyone!

5d ago 128 1

Arun Subramanian

Deploy Streamlit App on AWS EC2 instance

I've always been a firm believer in the philosophy of continuous learning. The worl...

Aug 31

See more recommendations